

## Red-Black Trees

### Balanced Insertions and Deletions Operations in $O(\lg n)$

11/15/07

RedBlack Trees

1

## Red-Black Trees

### ■ Def.: Red-Black Tree

A **red-black tree** is a binary search tree + 1 bit per node: an attribute *color*, which is either red or black.

All leaves are empty (*nil*) and colored black.

- We use a single sentinel,  $nil[T]$ , for all the leaves of red-black tree  $T$ .
- $color[nil[T]]$  is black.
- The root's parent is also  $nil[T]$ .

All other attributes of binary search trees are inherited by red-black trees (*key*, *left*, *right*, and *p*). We don't care about the key in  $nil[T]$ .

11/15/07

RedBlack Trees

2

## Red-Black Trees

### Red-black properties

1. Every node is either red or black.
2. The root is black.
3. Every leaf ( $nil[T]$ ) is black.
4. If a node is red, then both its children are black. (Hence no two reds in a row on a simple path from the root to a leaf.)
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes.


11/15/07

RedBlack Trees

3

## Red-Black Trees

### ■ Example:



11/15/07

RedBlack Trees


4

## Red-Black Trees

### ■ Height

Height of a red-black tree

- **Height of a node** is the number of edges in a longest path to a leaf.
- **Black-height** of a node  $x$ :  $bh(x)$  is the number of black nodes (including  $nil[T]$ ) on the path from  $x$  to leaf, not counting  $x$ . By property 5, black-height is well defined.



11/15/07

RedBlack Trees

5

## Red-Black Trees

### ■ Two sub-lemmata

**Claim**

Any node with height  $h$  has black-height  $\geq h/2$ .

**Proof** By property 4,  $\leq h/2$  nodes on the path from the node to a leaf are red. Hence  $\geq h/2$  are black. ■ (claim)

**Claim**

The subtree rooted at any node  $x$  contains  $\geq 2^{bh(x)} - 1$  internal nodes.

**Proof** By induction on height of  $x$ .

**Basis:** Height of  $x = 0 \Rightarrow x$  is a leaf  $\Rightarrow bh(x) = 0$ . The subtree rooted at  $x$  has 0 internal nodes.  $2^0 - 1 = 0$ .

**Inductive step:** Let the height of  $x$  be  $h$  and  $bh(x) = b$ . Any child of  $x$  has height  $h - 1$  and black-height either  $b$  (if the child is red) or  $b - 1$  (if the child is black). By the inductive hypothesis, each child has  $\geq 2^{bh(x)-1} - 1$  internal nodes. Thus, the subtree rooted at  $x$  contains  $\geq 2 \cdot (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$  internal nodes. (The +1 is for  $x$  itself.) ■ (claim)

11/15/07

RedBlack Trees

6

## Red-Black Trees

### ■ Another one

**Lemma**

A red-black tree with  $n$  internal nodes has height  $\leq 2 \lg(n + 1)$ .

**Proof** Let  $h$  and  $b$  be the height and black-height of the root, respectively. By the above two claims,

$$n \geq 2^b - 1 \geq 2^{h/2} - 1.$$

Adding 1 to both sides and then taking logs gives  $\lg(n + 1) \geq h/2$ , which implies that  $h \leq 2 \lg(n + 1)$ . ■ (theorem)

11/15/07

RedBlack Trees

7

## Red-Black Trees

### ■ Insertion & Deletion

If we insert, what color to make the new node?

- Red? Might violate property 4.
- Black? Might violate property 5.

If we delete, thus removing a node, what color was the node that was removed?

- Red? OK, since we won't have changed any black-heights, nor will we have created two red nodes in a row. Also, cannot cause a violation of property 2, since if the removed node was red, it could not have been the root.
- Black? Could cause there to be two reds in a row (violating property 4), and can also cause a violation of property 5. Could also cause a violation of property 2, if the removed node was the root and its child—which becomes the new root—was red.

11/15/07


RedBlack Trees

8

## Red-Black Trees

### Rotations

- Won't upset the binary-search-tree property.
- Have both left rotation and right rotation. They are inverses of each other.
- A rotation takes a red-black-tree and a node within the tree.



11/15/07

RedBlack Trees

9

## Red-Black Trees

### The Rotation Algorithm in Pseudo-Code

```

LEFT-ROTATE( $T, x$ )
 $y \leftarrow \text{right}[x]$            ▷ Set  $y$ .
 $\text{right}[x] \leftarrow \text{left}[y]$    ▷ Turn  $y$ 's left subtree into  $x$ 's right subtree.
if  $\text{left}[y] \neq \text{nil}[T]$ 
    then  $p[\text{left}[y]] \leftarrow x$ 
 $p[y] \leftarrow p[x]$            ▷ Link  $x$ 's parent to  $y$ .
if  $p[x] = \text{nil}[T]$ 
    then  $\text{root}[T] \leftarrow y$ 
else if  $x = \text{left}[p[x]]$ 
    then  $\text{left}[p[x]] \leftarrow y$ 
    else  $\text{right}[p[x]] \leftarrow y$ 
 $\text{left}[y] \leftarrow x$            ▷ Put  $x$  on  $y$ 's left.
 $p[x] \leftarrow y$ 
    
```

The pseudocode for LEFT-ROTATE assumes that

- $\text{right}[x] \neq \text{nil}[T]$ , and
- root's parent is  $\text{nil}[T]$ .

Pseudocode for RIGHT-ROTATE is symmetric: exchange *left* and *right* everywhere.


11/15/07

RedBlack Trees

10

## Red-Black Trees

### Rotation Example:



- Before rotation: keys of  $x$ 's left subtree  $\leq 11 \leq$  keys of  $y$ 's left subtree  $\leq 18 \leq$  keys of  $y$ 's right subtree.
- Rotation makes  $y$ 's left subtree into  $x$ 's right subtree.
- After rotation: keys of  $x$ 's left subtree  $\leq 11 \leq$  keys of  $x$ 's right subtree  $\leq 18 \leq$  keys of  $y$ 's right subtree.

**Time:**  $O(1)$  for both LEFT-ROTATE and RIGHT-ROTATE, since a constant number of pointers are modified.

11/15/07

RedBlack Trees


11

## Red-Black Trees

### Insertion Pseudo-Code

```

RB-INSERT( $T, z$ )
 $y \leftarrow \text{nil}[T]$ 
 $x \leftarrow \text{root}[T]$ 
while  $x \neq \text{nil}[T]$ 
    do  $y \leftarrow x$ 
       if  $\text{key}[z] < \text{key}[x]$ 
           then  $x \leftarrow \text{left}[x]$ 
           else  $x \leftarrow \text{right}[x]$ 
 $p[z] \leftarrow y$ 
if  $y = \text{nil}[T]$ 
    then  $\text{root}[T] \leftarrow z$ 
else if  $\text{key}[z] < \text{key}[y]$ 
    then  $\text{left}[y] \leftarrow z$ 
    else  $\text{right}[y] \leftarrow z$ 
 $\text{left}[z] \leftarrow \text{nil}[T]$ 
 $\text{right}[z] \leftarrow \text{nil}[T]$ 
 $\text{color}[z] \leftarrow \text{RED}$ 
RB-INSERT-FIXUP( $T, z$ )
    
```



11/15/07

RedBlack Trees

12

## Red-Black Trees

### ■ Insertion

- RB-INSERT ends by coloring the new node  $z$  red.
- Then it calls RB-INSERT-FIXUP because we could have violated a red-black property.

Which property might be violated?

1. OK.
2. If  $z$  is the root, then there's a violation. Otherwise, OK.
3. OK.
4. If  $p[z]$  is red, there's a violation: both  $z$  and  $p[z]$  are red.
5. OK.

Remove the violation by calling RB-INSERT-FIXUP:

11/15/07

RedBlack Trees

13

## Red-Black Trees

### ■ Fix the Problems

```

RB-INSERT-FIXUP( $T, z$ )
while  $color[p[z]] = RED$ 
do if  $p[z] = left[p[p[z]]]$ 
then  $y \leftarrow right[p[p[z]]]$ 
if  $color[y] = RED$ 
then  $color[p[z]] \leftarrow BLACK$ 
 $color[y] \leftarrow BLACK$ 
 $color[p[p[z]]] \leftarrow RED$ 
 $z \leftarrow p[p[z]]$ 
else if  $z = right[p[z]]$ 
then  $z \leftarrow p[z]$ 
LEFT-ROTATE( $T, z$ )
 $color[p[z]] \leftarrow BLACK$ 
 $color[p[p[z]]] \leftarrow RED$ 
RIGHT-ROTATE( $T, p[p[z]]$ )
else (same as then clause
with "right" and "left" exchanged)
 $color[root[T]] \leftarrow BLACK$ 
    
```

- ▷ Case 1
- ▷ Case 1
- ▷ Case 1
- ▷ Case 1
- ▷ Case 2
- ▷ Case 2
- ▷ Case 3
- ▷ Case 3
- ▷ Case 3

11/15/07

RedBlack Trees


14

## Red-Black Trees


### ■ Fix the Problems

```

RB-INSERT-FIXUP( $T, z$ )
while  $color[p[z]] = RED$ 
do if  $p[z] = left[p[p[z]]]$ 
then  $y \leftarrow right[p[p[z]]]$ 
if  $color[y] = RED$ 
then  $color[p[z]] \leftarrow BLACK$ 
 $color[y] \leftarrow BLACK$ 
 $color[p[p[z]]] \leftarrow RED$ 
 $z \leftarrow p[p[z]]$ 
else if  $z = right[p[z]]$ 
then  $z \leftarrow p[z]$ 
LEFT-ROTATE( $T, z$ )
 $color[p[z]] \leftarrow BLACK$ 
 $color[p[p[z]]] \leftarrow RED$ 
RIGHT-ROTATE( $T, p[p[z]]$ )
else (same as then clause
with "right" and "left" exchanged)
 $color[root[T]] \leftarrow BLACK$ 
    
```



$color[z] = red$



$color[z] = black$

11/15/07

RedBlack Trees


15

## Red-Black Trees

### ■ Insert another

```

RB-INSERT( $T, z$ )
 $y \leftarrow nil[T]$ 
 $x \leftarrow root[T]$ 
while  $x \neq nil[T]$ 
do  $y \leftarrow x$ 
if  $key[z] < key[x]$ 
then  $x \leftarrow left[x]$ 
else  $x \leftarrow right[x]$ 
 $p[z] \leftarrow y$ 
if  $y = nil[T]$ 
then  $root[T] \leftarrow z$ 
else if  $key[z] < key[y]$ 
then  $left[y] \leftarrow z$ 
else  $right[y] \leftarrow z$ 
 $left[z] \leftarrow nil[T]$ 
 $right[z] \leftarrow nil[T]$ 
 $color[z] \leftarrow RED$ 
RB-INSERT-FIXUP( $T, z$ )
    
```



11/15/07

RedBlack Trees


16

## Red-Back Trees

### ■ Fix it

```

RB-INSERT-FIXUP(T, z)
while color[p[z]] = RED
do if p[z] = left[p[p[z]]]
    then y ← right[p[p[z]]]
    if color[y] = RED
    then color[p[z]] ← BLACK
        color[y] ← BLACK
        color[p[p[z]]] ← RED
        z ← p[p[z]]
    else if z = right[p[z]]
    then z ← p[z]
        LEFT-ROTATE(T, z)
        color[p[z]] ← BLACK
        color[p[p[z]]] ← RED
        RIGHT-ROTATE(T, p[p[z]])
    else (same as then clause
        with "right" and "left" exchanged)
color[root[T]] ← BLACK
    
```



Nothing to do - root already black

11/15/07

RedBlack Trees


17

## Red-Black Trees

### ■ Insert another

```

RB-INSERT(T, z)
y ← nil[T]
x ← root[T]
while x ≠ nil[T]
do y ← x
    if key[z] < key[x]
    then x ← left[x]
    else x ← right[x]
p[z] ← y
if y = nil[T]
then root[T] ← z
else if key[z] < key[y]
then left[y] ← z
    else right[y] ← z
left[z] ← nil[T]
right[z] ← nil[T]
color[z] ← RED
RB-INSERT-FIXUP(T, z)
    
```



11/15/07

RedBlack Trees


18

## Red-Black Trees

### ■ And fix it

```

RB-INSERT-FIXUP(T, z)
while color[p[z]] = RED
do if p[z] = left[p[p[z]]]
    then y ← right[p[p[z]]]
    if color[y] = RED
    then color[p[z]] ← BLACK
        color[y] ← BLACK
        color[p[p[z]]] ← RED
        z ← p[p[z]]
    else if z = right[p[z]]
    then z ← p[z]
        LEFT-ROTATE(T, z)
        color[p[z]] ← BLACK
        color[p[p[z]]] ← RED
        RIGHT-ROTATE(T, p[p[z]])
    else (same as then clause
        with "right" and "left" exchanged)
color[root[T]] ← BLACK
    
```



11/15/07

RedBlack Trees


19

## Red-Black Trees

### ■ And fix it

```

else /* if p[z] = right[p[p[z]]] then */
y ← left[p[p[z]]]
if color[y] = RED
then color[p[z]] ← BLACK
    color[y] ← BLACK
    color[p[p[z]]] ← RED
    z ← p[p[z]]
else if z = left[p[z]]
then z ← p[z]
    Right-Rotate(T, z)
color[p[z]] ← BLACK
color[p[p[z]]] ← RED
Left-Rotate(T, p[p[z]])
color[root[T]] ← BLACK
    
```



Note: y is BLACK and z is NOT a left child of its parent, so we color p[z] BLACK, p[p[z]] RED and Left-Rotate on p[p[z]] and finish by coloring p[z] (the new root) BLACK.


11/15/07

RedBlack Trees

20

## Red-Black Trees

- Here is the sequence: color  $p[z]$  BLACK,  $p[p[z]]$  RED and Left-Rotate on  $p[p[z]]$  and finish by coloring  $p[z]$  (the new root) BLACK.




11/15/07

RedBlack Trees

21

## Red-Black Trees

- Insert e. The problem, at this point, is that the number of black nodes along each path must change. Look back at the code: does this apply?



```

RB-INSERT-FIXUP(T, z)
while color[p[z]] = RED
do if p[z] = left[p[p[z]]]
then y ← right[p[p[z]]]
if color[y] = RED
then color[p[z]] ← BLACK
color[y] ← BLACK
color[p[p[z]]] ← RED
z ← p[p[z]]
else if z = right[p[p[z]]]
then z ← p[p[z]]
LEFT-ROTATE(T, z)
color[p[z]] ← BLACK
color[p[p[z]]] ← RED
RIGHT-ROTATE(T, p[p[z]])
else (same as then clause
with "right" and "left" exchanged)
color[root[T]] ← BLACK
    
```


11/15/07

RedBlack Trees

22

## Red-Black Trees

- Look back at the code: does this apply?



```

else /* if p[z] = right[p[p[z]]] then */
y ← left[p[p[z]]]
if color[y] = RED
then color[p[z]] ← BLACK
color[y] ← BLACK
color[p[p[z]]] ← RED
z ← p[p[z]]
else if z = left[p[p[z]]]
then z ← p[p[z]]
Right-Rotate(T, z)
color[p[z]] ← BLACK
color[p[p[z]]] ← RED
Left-Rotate(T, p[p[z]])
color[root[T]] ← BLACK
    
```

Since y is RED, set  $color[p[z]]$  to BLACK,  $color[y]$  to BLACK,  $color[p[p[z]]]$  to RED,  $z$  to  $p[p[z]]$ ; since now  $color[p[z]] = color[p[root[T]]] = BLACK$ , the while loop ends. Set  $z = root[T]$  to BLACK.


11/15/07

RedBlack Trees

23

## Red-Black Trees

- And the tree looks like:




11/15/07

RedBlack Trees

24

## Red-Black Trees

- Add f: Do we use the **left** or **right** code? We start using the **left** code.



```

RB-INSERT-FIXUP(T, z)
while color[p[z]] = RED
do if p[z] = left[p[p[z]]]
then y ← right[p[p[z]]]
if color[y] = RED
then color[p[z]] ← BLACK
color[y] ← BLACK
color[p[p[z]]] ← RED
z ← p[p[z]]
else if z = right[p[p[z]]]
then z ← p[z]
LEFT-ROTATE(T, z)
color[p[z]] ← BLACK
color[p[p[z]]] ← RED
RIGHT-ROTATE(T, p[p[z]])
else (same as then clause
with "right" and "left" exchanged)
color[root[T]] ← BLACK
    
```

11/15/07

RedBlack Trees

25

## Red-Black Trees

- Why does this work?

### Loop invariant:

At the start of each iteration of the **while** loop,

- z is red.
- There is at most one red-black violation:
  - Property 2: z is a red root, or
  - Property 4: z and p[z] are both red.

11/15/07

RedBlack Trees

26

## Red-Black Trees

- The Induction

**Initialization:** We've already seen why the loop invariant holds initially.

**Termination:** The loop terminates because p[z] is black. Hence, property 4 is OK. Only property 2 might be violated, and the last line fixes it.

**Maintenance:** We drop out when z is the root (since then p[z] is the sentinel nil[T], which is black). When we start the loop body, the only violation is of property 4.

There are 6 cases, 3 of which are symmetric to the other 3. The cases are not mutually exclusive. We'll consider cases in which p[z] is a left child.

Let y be z's uncle (p[z]'s sibling).

11/15/07


RedBlack Trees

27

## Red-Black Trees

- Case 1

Case 1: y is red



If z is a right child

If z is a left child

- p[p[z]] (z's grandparent) must be black, since z and p[z] are both red and there are no other violations of property 4.
- Make p[z] and y black ⇒ now z and p[z] are not both red. But property 5 might now be violated.
- Make p[p[z]] red ⇒ restores property 5.
- The next iteration has p[p[z]] as the new z (i.e., z moves up 2 levels).

11/15/07


RedBlack Trees

28

## Red-Black Trees

### Case 2

Case 2:  $y$  is black,  $z$  is a right child



- Left rotate around  $p[z] \Rightarrow$  now  $z$  is a left child, and both  $z$  and  $p[z]$  are red.
- Takes us immediately to case 3.

11/15/07

RedBlack Trees

29

## Red-Black Trees

### Case 3

Case 3:  $y$  is black,  $z$  is a left child

- Make  $p[z]$  black and  $p[p[z]]$  red.
- Then right rotate on  $p[p[z]]$ .
- No longer have 2 reds in a row.
- $p[z]$  is now black  $\Rightarrow$  no more iterations.


11/15/07

RedBlack Trees

30

## Red-Black Trees

### Finish the last insertion (modified).




11/15/07

RedBlack Trees

31

## Red-Black Trees

### Now add g:



11/15/07

RedBlack Trees

32




## Red-Black Trees

### ■ Now for the deletion: delete g.

```

RB-DELETE( $T, z$ )
if  $left[z] = nil[T]$  or  $right[z] = nil[T]$ 
  then  $y \leftarrow z$ 
  else  $y \leftarrow \text{TREE-SUCCESSOR}(z)$ 
if  $left[y] \neq nil[T]$ 
  then  $x \leftarrow left[y]$ 
  else  $x \leftarrow right[y]$ 
 $p[x] \leftarrow p[y]$ 
if  $p[y] = nil[T]$ 
  then  $root[T] \leftarrow x$ 
  else if  $y = left[p[y]]$ 
    then  $left[p[y]] \leftarrow x$ 
    else  $right[p[y]] \leftarrow x$ 
if  $y \neq z$ 
  then  $key[z] \leftarrow key[y]$ 
  copy  $y$ 's satellite data into  $z$ 
if  $color[y] = \text{BLACK}$ 
  then RB-DELETE-FIXUP( $T, x$ )
return  $y$ 

```



dBlack Trees

33

## Red-Black Trees

- Chasing the code with the picture (or vice-versa, your pick), we end up just removing the node labeled g. Nothing else changes, other than the left pointer out of the node labeled h, which goes to  $nil[T]$ . g just disappears without triggering any adjustments: it was a RED node, so the number of BLACK nodes along that path did not change. Notice that no labels or other data are copied, since the node we are deleting is the very last in a chain - the next nodes are the sentinel leaves.
- Deleting any other node will trigger more complicated readjustments.

11/15/07

RedBlack Trees

34


## Red-Black Trees

### ■ Remove f:

```

RB-DELETE( $T, z$ )
if  $left[z] = nil[T]$  or  $right[z] = nil[T]$ 
  then  $y \leftarrow z$ 
  else  $y \leftarrow \text{TREE-SUCCESSOR}(z)$ 
if  $left[y] \neq nil[T]$ 
  then  $x \leftarrow left[y]$ 
  else  $x \leftarrow right[y]$ 
 $p[x] \leftarrow p[y]$ 
if  $p[y] = nil[T]$ 
  then  $root[T] \leftarrow x$ 
  else if  $y = left[p[y]]$ 
    then  $left[p[y]] \leftarrow x$ 
    else  $right[p[y]] \leftarrow x$ 
if  $y \neq z$ 
  then  $key[z] \leftarrow key[y]$ 
  copy  $y$ 's satellite data into  $z$ 
if  $color[y] = \text{BLACK}$ 
  then RB-DELETE-FIXUP( $T, x$ )
return  $y$ 

```



dBlack Trees

35

## Red-Black Trees

- This is also easy, since removing f involves finding its successor (g), re-attaching the parent of g to a sentinel (rather than g), and copying the contents of g into f. Since the node actually removed (g) is RED, nothing needs to be done.
- We now try to remove h - this is a BLACK node and, because it has only one child, it will actually be removed. This will finally trigger RB-Delete-Fixup( $T, x$ ).


11/15/07

RedBlack Trees

36

## Red-Black Trees

- Before the call to RB-Delete-Fixup( $T, x$ ) We have:



11/15/07

RedBlack Trees

37

## Red-Black Trees

- Here is the pseudo-code

```

RB-DELETE-FIXUP( $T, x$ )
while  $x \neq \text{root}[T]$  and  $\text{color}[x] = \text{BLACK}$ 
do if  $x = \text{left}[p[x]]$ 
then  $w \leftarrow \text{right}[p[x]]$ 
if  $\text{color}[w] = \text{RED}$ 
then  $\text{color}[w] \leftarrow \text{BLACK}$            > Case 1
    $\text{color}[p[x]] \leftarrow \text{RED}$        > Case 1
   LEFT-ROTATE( $T, p[x]$ )           > Case 1
    $w \leftarrow \text{right}[p[x]]$        > Case 1
if  $\text{color}[\text{left}[w]] = \text{BLACK}$  and  $\text{color}[\text{right}[w]] = \text{BLACK}$ 
then  $\text{color}[w] \leftarrow \text{RED}$        > Case 2
    $x \leftarrow p[x]$                > Case 2
else if  $\text{color}[\text{right}[w]] = \text{BLACK}$ 
then  $\text{color}[\text{left}[w]] \leftarrow \text{BLACK}$  > Case 3
    $\text{color}[w] \leftarrow \text{RED}$        > Case 3
   RIGHT-ROTATE( $T, w$ )           > Case 3
    $w \leftarrow \text{right}[p[x]]$        > Case 3
    $\text{color}[w] \leftarrow \text{color}[p[x]]$  > Case 4
    $\text{color}[p[x]] \leftarrow \text{BLACK}$    > Case 4
    $\text{color}[\text{right}[w]] \leftarrow \text{BLACK}$  > Case 4
   LEFT-ROTATE( $T, p[x]$ )         > Case 4
    $x \leftarrow \text{root}[T]$ 
else (same as then clause with "right" and "left" exchanged)
 $\text{color}[x] \leftarrow \text{BLACK}$ 
    
```


11/15/07

RedBlack Trees

38

## Red-Black Trees

- Since  $x$  is neither the root, nor is it BLACK (it took the place of a BLACK node that was removed), the code tells us to just color it BLACK.



11/15/07


RedBlack Trees

39

## Red-Black Trees

- How about removing  $e$  from the original tree? Since the deletion itself does not worry about color, we just remove  $e$ , and  $x$  is the  $\text{nil}[T]$  sentinel. Note that the parent of the sentinel is now  $f$ .

Furthermore,  $x$  is not the root, and its color is BLACK.



11/15/07

RedBlack Trees


40

## Red-Black Trees

### ■ Compare the code and the tree:

```

RB-DELETE-FIXUP(T, x)
while x ≠ root[T] and color[x] = BLACK
do if x = left[p[x]]
    then w ← right[p[x]]
    if color[w] = RED
    then color[w] ← BLACK
    color[p[x]] ← RED
    LEFT-ROTATE(T, p[x])
    w ← right[p[x]]
    if color[left[w]] = BLACK and color[right[w]] = BLACK
    then color[w] ← RED
    x ← p[x]
    else if color[right[w]] = BLACK
    then color[left[w]] ← BLACK
    color[w] ← RED
    RIGHT-ROTATE(T, w)
    w ← right[p[x]]
    color[w] ← color[p[x]]
    color[p[x]] ← BLACK
    color[right[w]] ← BLACK
    LEFT-ROTATE(T, p[x])
    x ← root[T]
else (same as then clause with "right" and "left" exchanged)
color[x] ← BLACK
    
```



11/15/07

RedBlack Trees

41

## Red-Black Trees

- We have a number of cases to take care of - 8 (precisely but not mutually exclusive). We will look at 4 of them, leaving the 4 symmetric ones as exercises.
- We observe that splicing out ANY RED node requires no fix-up: it does not alter the black height of any node; it does not introduce any pair of parent-child RED nodes; and it does not change the root.
- The only case where fixing up will be needed is when the spliced out node is BLACK: that will alter the black height of its ancestors, and may violate the requirement that there be no pair of parent-child RED nodes.

11/15/07


RedBlack Trees

42

## Red-Black Trees

How can the splicing out of a black node  $y$  effect the result?

1.  $y$  had been the root and a red child of  $y$  becomes (physically) the new root, violating property 2. This can occur only if we have the configuration on the left, or its symmetric counterpart.



The fix is to just color the new root BLACK, but we will examine it in the context of RB-Delete-Fixup.

11/15/07

RedBlack Trees

43

## Red-Black Trees

2. Both  $x$  and  $p[x] = p[y]$  are RED. Property 4 is violated (red has only black children).
3. If  $y$  was BLACK, its removal from any path will cause any path that contained it to have one fewer BLACK nodes. Property 5 (all paths from node down have same number of black nodes) is violated by any ancestor of  $y$  in the tree.

How do we solve the problem? pretend that the node  $x$  has an extra BLACK. Then all is well (properties 4 & 5), and we have to figure out where to unload this extra black...


11/15/07

RedBlack Trees

44

## Red-Black Trees

- Note that, in  $\text{RB-Delete-Fixup}(T, x)$ , with  $y$  the node actually deleted,  $x$  is
  - a.  $y$ 's sole non-sentinel child before  $y$  was spliced out
  - b. the sentinel itself, if  $y$  had no children.



Also, after deletion,  $p[x] = p[y]$

11/15/07

RedBlack Trees

45

## Red-Black Trees

With  $y$  black, we could violate the properties:

1. Every node is either red or black. NO
2. The root is black: YES, if  $y$  is the root and  $x$  is red.
3. Every leaf is black. NO
4. If a node is red then both of its children are black: YES, if  $p[y]$  and  $x$  are both red (from the left hand example).
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes: YES, any path that contained  $y$  now has one fewer black nodes.

We can fix 5 by giving  $x$  "an extra black" from its **deleted parent** - then the count of black nodes if "fixed": we will push the extra black around until we can safely unload it...

11/15/07

RedBlack Trees

46

## Red-Black Trees

Note: we have violated Property 1, since we now have nodes that are neither red nor black.:  $x$  is doubly-black if  $x$  was black; it is red&black if it was red. Note that  $\text{color}[x]$  is still just RED or BLACK - the extra black comes from **pointing** to it (which means we have to unload its extra blackness **before** we stop pointing to it).

**IDEA:** move the extra black up the tree until:

- $x$  points to a red&black node --> turn it into a red one
- $x$  points to the root --> just remove the extra black
- perform rotations and recolorings

The first two points tell you when it's safe to unload the extra black; the last one tells you how to move it up. That's where we go now.

11/15/07


RedBlack Trees

47

## Red-Black Trees

We start by assuming  $x$  is the left child of its parent and that  $w$  is its sibling:

Case 1:  $w$  is red



- $w$  must have black children.
- Make  $w$  black and  $p[x]$  red.
- Then left rotate on  $p[x]$ .
- New sibling of  $x$  was a child of  $w$  before rotation  $\Rightarrow$  must be black.
- Go immediately to case 2, 3, or 4.

The same black path conditions will be satisfied at the end of the recoloring and rotation as at the beginning.

11/15/07


RedBlack Trees

48

## Red-Black Trees

The tree rooted at A has one more black node than the trees rooted at B and C. We don't know the color of the children of C - assume they are both black (other cases later): A, B, D are, respectively, the old A, B, C.

Case 2:  $w$  is black and both of  $w$ 's children are black



[Node with gray outline is of unknown color, denoted by  $c$ .]

- Take 1 black off  $x$  ( $\Rightarrow$  singly black) and off  $w$  ( $\Rightarrow$  red).
- Move that black to  $p[x]$ .
- Do the next iteration with  $p[x]$  as the new  $x$ .
- If entered this case from case 1, then  $p[x]$  was red  $\Rightarrow$  new  $x$  is red & black  $\Rightarrow$  color attribute of new  $x$  is RED  $\Rightarrow$  loop terminates. Then new  $x$  is made black in the last line.

11/15/07


RedBlack Trees

49

## Red-Black Trees

Next case: doesn't quite end, since we can't yet get rid of the extra black on A. You just move to the case where the colors of the children of  $w$  are swapped ( $\gamma$  is black)

Case 3:  $w$  is black,  $w$ 's left child is red, and  $w$ 's right child is black



- Make  $w$  red and  $w$ 's left child black.
- Then right rotate on  $w$ .
- New sibling  $w$  of  $x$  is black with a red right child  $\Rightarrow$  case 4.

11/15/07


RedBlack Trees

50

## Red-Black Trees

Finally:

Case 4:  $w$  is black,  $w$ 's left child is black, and  $w$ 's right child is red



[Now there are two nodes of unknown colors, denoted by  $c$  and  $c'$ .]

- Make  $w$  be  $p[x]$ 's color ( $c$ ).
- Make  $p[x]$  black and  $w$ 's right child black.
- Then left rotate on  $p[x]$ .
- Remove the extra black on  $x$  ( $\Rightarrow$   $x$  is now singly black) without violating any red-black properties.
- All done. Setting  $x$  to root causes the loop to terminate.

11/15/07

RedBlack Trees

51

## Red-Black Trees

11/15/07

RedBlack Trees

52

## Red-Black Trees

---

11/15/07

RedBlack Trees

53