

Solutions to First Examination

CS 430 Introduction to Algorithms
Spring, 2017

Monday, February 20, 2017
10am–11:15am, 002 Herman Hall
11:25am–12:40pm, 104 Rettaliata Engineering Center

Exam Statistics

130 students took the exam; there were 3 no-shows recorded as 0. The range of scores was 0–90, with a mean of 44.76 (this includes the no-shows), a median of 48, and a standard deviation of 18.05. Very roughly speaking, if I had to assign final grades on the basis of this exam only, above 60 would be an A (25), 46–59 a B (46), 30–45 a C (34), 20–29 a D (16), below 20 an E (12). Every student should have been able to get full credit, or nearly so, on the first, fourth and last problems, plus a few points on the other two problems; that is, no score should have been much below 60.

Problem Solutions

- (a) $T(n)$ grows linearly because the annihilator is $(E - 1)^2$ so $T(2n) = 2T(n)$; hence the running time doubles.
 - (b) $T(n)$ grows cubically because the annihilator is $(E - 1)^4$ so $T(2n) = 8T(n)$; hence the running time is multiplied by 8.
 - (c) $T(n)$ grows like the Fibonacci numbers; $F_n = \Theta(\phi^n)$ where $\phi = (1 + \sqrt{5})/2 \approx 1.61801$ so $T(2n) = \Theta(\phi^{2n}) = \Theta(T(n)^2)$. Hence the running time is roughly squared.
 - (d) The recurrence tells us that $T(2n) = T(n) + 1/(2n)$ so the running time is only negligibly affected.
- Let the elements of list S (ordered by starting value) be $[a_i, b_i]$ and let the elements of list E (ordered by ending value) be $[\hat{a}_i, \hat{b}_i]$. We merge S and E looking at the starting values a_i of S and the ending values \hat{b}_i of E into a list of $2n$ values—think of the a_i as opening parentheses and the \hat{b}_i as closing parentheses. We start a counter at 0 and, as the merge is being done, we increment the counter when an a_i is put into the merged output and decrease the counter when a \hat{b}_i is put into the merged output. Of course the counter is 0 at the end (why?). The maximum value attained by the counter over the course of the merge is the maximum nesting depth of the intervals.
- (a) If n is even, then because the elements are in decreasing order the element 1 is in an even position; the restricted swaps mean that it can only be moved to other even positions, never to the first position as needed to sort the elements.

- (b) Suppose $n = 2k + 1$ so that $k = (n - 1)/2$. Use insertion sort separately on the $k + 1$ odd positions $1, 3, \dots, 2k + 1$ and then on the k even positions $2, 4, \dots, 2k$; the resulting array is now fully sorted.
- (c) In the worst case every comparison results in a swap. As we saw in the notes for Lecture 3 (January 18), the worst case for insertion sort to sort the $k + 1$ odd-indexed values takes $(k + 1)(k)/2$ comparisons; the worst case to sort the k even-indexed values takes $(k)(k - 1)/2$ comparisons. In total then $k^2 = [(n - 1)/2]^2 = (n - 1)^2/4$ comparisons (and hence swaps) are used in the worst case.
- (d) Following the hint, we count the number of inversions in the given input. The $k + 1$ odd-indexed elements are in reverse order, and hence have $0 + 1 + 2 + \dots + k = (k)(k + 1)/2 = k(k + 1)/2$ inversions. Similarly, the k even-indexed elements have $0 + 1 + 2 + \dots + (k - 1) = k(k - 1)/2$ inversions. There are thus $k(k + 1)/2 + k(k - 1)/2 = k^2$ inversions, each of which must be undone by a swap; hence $k^2 = (n - 1)^2/4$ swaps are needed in the worst case.
4. For $n = 1$ the tree has a single leaf at depth 0, so $\sum_{i=1}^n 2^{-l_i} = 2^0 = 1$. Now suppose the inequality is true for all tree with fewer than n leaves; that means that in a tree of n leaves both the left and right subtrees satisfy the inequality so

$$\sum_{\text{leaves } l \text{ in left subtree}} 2^{-(\text{depth}(l)-1)} \leq 1$$

and

$$\sum_{\text{leaves } l \text{ in right subtree}} 2^{-(\text{depth}(l)-1)} \leq 1$$

because the leaves are one level shallower with respect to the roots of the subtrees. Adding these two inequalities and dividing by 2 give the desired result,

$$\sum_{\text{leaves } l \text{ in tree}} 2^{-\text{depth}(l)} \leq 1$$

5. We know from Lemma 1 of Lecture 6 (January 30) that the binary tree T with the least external path length has all of its leaves on levels l and $l + 1$ for some value of l . So coloring all internal nodes at level l red and all other internal nodes black gives every leaf a black depth of l . There are clearly no two parent/child red nodes and the root is black. Hence this is a proper red-black coloring of the tree.