

CS 553: Cloud Computing - HA1

Saptarshi Chatterjee
CWID: A20413922

March 28, 2018

1 Hardware Specification

Experiments are run on 2 different clusters -

- **Hyperion** : Intel Xeon E312xx (Sandy Bridge). MemTotal: 4046452 kB
. Ethernet interface product: Virtio network device
- **Prometheus** Disk- is Micron 5100 PRO 2.5" 480GB, SATA, 6Gb/s, 3D NAND, 7mm, 1.5D WPD2,
NIC- Super Micro Computer Inc Ethernet Controller 10-Gigabit X540-AT2, CPU - Intel(R) Xeon(R) CPU E5-2620 v4 @ 2.10GHz , 8GiB DIMM
Synchronous 2400 MHz (0.4 ns, 64bit)

2 CPU Benchmarking (Run on Hyperion)

We were asked to perform 1 trillion arithmetic (quarter precision, half precision, single precision, double precision) using 1, 2 and 4 threads.

$$\begin{aligned}\text{FLOPS} &= \text{sockets} \times \frac{\text{cores}}{\text{socket}} \times \frac{\text{cycles}}{\text{second}} \times \frac{\text{FLOPs}}{\text{cycle}} \\ &= 2 * 1 * 2.299 * 8 \\ &= 36.78(\text{For DP})(\text{Sandy Bridge -8 DP FLOPs/cycle}) \\ &= 73.56(\text{for SP})(16 \text{ SP FLOPs/cycle: 8-wide AVX addition} + 8\text{-wide AVX multiplication}) \\ &= 147.136(\text{for QP}) \\ &= 147.136(\text{for HP})\end{aligned}$$

[illegible]

```

Decisions: DP, Threads 4, TotalTime = 58.147661, Gflops: 17.497637
schrotterjee@lacompute-7:~$ ./reports/home/schrotterjee/l_mklbp_2018.1.030/benchmarks_2018/linux/mkl/benchmarks/linpack/2d/linpack_xeor64 /exports/home/schrotterjee/l_mklbp_2018.1.030/benchmarks_2018/linux/mkl/benchmarks/linpack/linpack_xeor64
Sample data file lininput_xeor64.

Current date/time: Sun Mar 25 04:51:25 2018

CPU frequency: 2.754 Gz
Number of CPUs: 2
Number of cores: 2
Number of threads: 2

Parameters are set to:

Number of tests: 15
Number of equations to solve (problem size) : 1030 2030 3030 1030 1530 1830 2230 2530 3030 2703 3303 5503 6303 6503 6703
Number of dimensions in matrix : 1030 2030 3030 1030 1530 1830 2230 2530 3030 2703 3303 5503 6303 6503 6703
Number of trials to run : 4 2 2 2 2 2 2 2 2 2 1 1 1 1 1
Data alignment value (in bytes) : 4 4 4 4 4 4 4 4 4 4 1 1 1 1 1

Maximum memory requested that can be used:335594416, at the size:20300

=====
Timing 1 linear equation system solver
=====
Size LDA Align Time(s) GFlops residual Residual(norm) Check
1030 1030 4 0.015 43.6410 0.394430e-13 3.237742e-02 pass
1030 1030 4 0.013 49.9520 0.394430e-13 3.237742e-02 pass
1030 1030 4 0.013 50.9627 0.394430e-13 3.237742e-02 pass
1030 1030 4 0.014 49.5592 0.394430e-13 3.237742e-02 pass
2030 2030 4 0.032 57.7435 4.085732e-12 3.554085e-02 pass
2030 2030 4 0.035 62.9229 4.085732e-12 3.554085e-02 pass
2030 2030 4 1.220 55.7530 2.205256e-11 3.454932e-02 pass
2030 5030 4 1.243 67.1011 2.205256e-11 3.454932e-02 pass
10303 10303 4 10.054 65.2055 0.157991e-11 3.239775e-02 pass
10303 10303 4 9.543 69.8012 0.157991e-11 3.239775e-02 pass

```

| Workload | Concurrency | MyCPUBench Measured Ops/Sec | HPL Measured Ops/Sec (GigaOPS) | Theoretical Ops/Sec (GigaOPS) | MyCPUBench Efficiency (%) | HPL Efficiency (%) |
|----------|-------------|--------------------------------|-----------------------------------|----------------------------------|------------------------------|-----------------------|
| QP | 1 | 45.393182 | N/A | 73.568 | 61.70234613 | N/A |
| QP | 2 | 90.078258 | N/A | 147.136 | 61.22108661 | N/A |
| QP | 4 | 85.05064 | N/A | 294.272 | 28.9020498 | N/A |
| HP | 1 | 46.200567 | N/A | 73.568 | 62.79981378 | N/A |
| HP | 2 | 86.669638 | N/A | 147.136 | 58.90444079 | N/A |
| HP | 4 | 85.924472 | N/A | 294.272 | 29.19899685 | N/A |
| SP | 1 | 23.579637 | N/A | 36.784 | 64.1029714 | N/A |
| SP | 2 | 46.652203 | N/A | 73.568 | 63.41371656 | N/A |
| SP | 4 | 43.583872 | N/A | 147.136 | 29.6214876 | N/A |
| DP | 1 | 8.508647 | 11.1738 | 18.392 | 46.26276098 | 60.75358852 |
| DP | 2 | 16.862083 | 24.9512 | 36.784 | 45.8408085 | 67.83166594 |
| DP | 4 | 17.197537 | 69.218 | 73.568 | 23.37638239 | 94.08710309 |

- Getting best result for QP with 2 threads . 90.07 GLOPS.
- GFLOPS is lowest for DP with 1 thread as we are not optimally using the available CPU power.
- I'm getting around 23% to 69% efficiency varying between different pre-

cision and concurrency using AVX. Where as Linpack gets 60%-94% efficiency

- It's not a Quad-core machine so using 4 threads actually hurts performance than 2 threads as a lot of time is spent in context switch .
- Linpack gets a best performance of 69.2180 GFLOPS for problem size 20000.

3 Memory Benchmarking (Run on Hyperion)

We were asked to perform random and Sequential Read/Write Ops on memory 1GB Workload. Operated over it 100X times with various access patterns (RWS, RWR) , various block sizes (1KB, 1MB, 10MB) and Concurrency (1 thread, 2 threads, 4 threads)

Theoretical Bandwidth = Base DRAM clock frequency \times Number of data transfers per clock
 \times Memory bus (interface) width \times Number of interfaces
 $= 2133,000,000 * 2 * 64 * 1 / (8 * 1000 * 1000 * 1000)$
 \dots Assuming clock frequency 2133MHz
 $= 34.128GBps$

Executing Memory benchmark

```
schatterjee@hyperionides: $ ./cs553-pat/memory_engine.sh
/exports/home/schatterjee/cs553-pat/memory/MRWBench.cpp: In function 'int main(int, char**)':
/exports/home/schatterjee/cs553-pat/memory/MRWBench.cpp:103:63: warning: ISO C++ forbids converting a string constant to 'char*' [-Wwrite-strings]
    error("ERROR in creating p_thread");
    ^
/exports/home/schatterjee/cs553-pat/memory/MRWBench.cpp:113:68: warning: ISO C++ forbids converting a string constant to 'char*' [-Wwrite-strings]
    error("ERROR in p_thread join");
    ^
/exports/home/schatterjee/cs553-pat/memory/MRWBench.cpp:129:55: warning: ISO C++ forbids converting a string constant to 'char*' [-Wwrite-strings]
    storeInFile(reportPath, report.c_str());
    ^
Folder /slurms and /reports created..
Running Memory tests
Submitted batch job 19188
schatterjee@hyperionides: $ squeue
JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)
12233 compute ocl ochen41 PD 0:00 1 (PartitionTimeLimit)
12214 compute ocl ochen41 PD 0:00 1 (PartitionTimeLimit)
13005 interact1 bash aveline R 34:05 1 bluecompute-1
18108 interact1 bash med3 R 25:00 1 bluecompute-2
18110 interact1 bash asing63 R 24:19 1 bluecompute-3
18109 compute NET-UDF8 iweeli R 9:32 2 redcompute-[20-21]
18170 compute disk-slu asing63 R 8:46 1 redcompute-18
18181 compute memory.s schatter R 7:56 1 redcompute-29
```

```
42954
Threads 2, BLOCKSIZE -1030, TotalTime - 53.016736, TotalData - GB:100, AccessType -RMW, Throughput - GFS:1.8
65410
Threads 2, BLOCKSIZE -1030, TotalTime - 32.846752, TotalData - GB:100, AccessType -RMS, Throughput - GFS:3.0
44230
Threads 1, BLOCKSIZE -1030, TotalTime - 96.555199, TotalData - GB:100, AccessType -RMW, Throughput - GFS:1.0
53339
Threads 1, BLOCKSIZE -1030, TotalTime - 62.539427, TotalData - GB:100, AccessType -RMS, Throughput - GFS:1.5
09445
===== Time: Sun Mar 25 12:24:01 CDT 2018. New Execution =====
Threads 4, BLOCKSIZE -1030330, TotalTime - 33.139563, TotalData - GB:100, AccessType -RMW, Throughput - GFS
13.617297
```

pmbw - Parallel Memory Bandwidth Benchmark / Measurement

corresponding results re kept in /cs553-pa1/memory/pmbw.txt

$$Latency(ns) = clockcycletime(ns) \times numberofclockcycles =$$

| Workload | Concurrency | Block Size | MyRAMBench Measured Throughput (GB/sec) | pmbw Measured Throughput (GB/sec) | Theoretical Throughput (GB/sec) | MyRAMBench Efficiency (%) | pmbw Efficiency (%) |
|----------|-------------|------------|--|--------------------------------------|------------------------------------|------------------------------|---------------------|
| RWS | 1 | 1 KB | 1.63132 | 17.3 | 34.128 | 4.78000469 | 50.6915143 |
| RWS | 1 | 1 MB | 1.69056 | 16.22 | 34.128 | 4.9535865 | 47.52695734 |
| RWS | 1 | 10MB | 1.616496 | 15.59 | 34.128 | 4.73656821 | 45.68096578 |
| RWS | 2 | 1 KB | 3.151194 | 29.77 | 34.128 | 9.2334564 | 87.23042663 |
| RWS | 2 | 1 MB | 3.269755 | 29.92 | 34.128 | 9.58085736 | 87.66994843 |
| RWS | 2 | 10MB | 2.97559 | 30.67 | 34.128 | 8.71891116 | 89.86755743 |
| RWS | 4 | 1 KB | 3.150432 | 26.42 | 34.128 | 9.23122363 | 77.41443976 |
| RWS | 4 | 1 MB | 3.234842 | 34.08 | 34.128 | 9.4785572 | 99.85935302 |
| RWS | 4 | 10MB | 3.008679 | 30.71 | 34.128 | 8.81586674 | 89.98476324 |
| RWR | 1 | 1 KB | 1.137239 | 17.3 | 34.128 | 3.33227555 | 50.6915143 |
| RWR | 1 | 1 MB | 1.53018 | 16.22 | 34.128 | 4.48364979 | 47.52695734 |
| RWR | 1 | 10MB | 1.566664 | 15.59 | 34.128 | 4.59055321 | 45.68096578 |
| RWR | 2 | 1 KB | 1.975065 | 29.77 | 34.128 | 5.7872275 | 87.23042663 |
| RWR | 2 | 1 MB | 2.928384 | 29.92 | 34.128 | 8.58059072 | 87.66994843 |
| RWR | 2 | 10MB | 3.030498 | 30.67 | 34.128 | 8.87979958 | 89.86755743 |
| RWR | 4 | 1 KB | 2.065833 | 26.42 | 34.128 | 6.05319093 | 77.41443976 |
| RWR | 4 | 1 MB | 2.978313 | 34.08 | 34.128 | 8.72688994 | 99.85935302 |
| RWR | 4 | 10MB | 3.017597 | 30.71 | 34.128 | 8.84199777 | 89.98476324 |

| Workload | Concurrency | Block Size | MyRAMBench Measured Latency (us) | pmbw Measured Latency (us) | Theoretical Latency (us) | MyRAMBench Efficiency (%) | pmbw Efficiency (%) |
|----------|-------------|------------|----------------------------------|----------------------------|--------------------------|---------------------------|---------------------|
| RWS | 1 1B | | 0.000613001 | 5.78035e-05 | 0.004 | 84.6798713 | 98.55491329 |
| RWS | 2 1B | | 0.00031734 | 3.35909e-05 | 0.004 | 92.06499492 | 99.160228484 |
| RWS | 4 1B | | 0.00017417 | 1.78035e-05 | 0.004 | 92.06458035 | 99.05374716 |
| RWR | 1 1B | | 0.000879323 | 5.78035e-05 | 0.004 | 78.016934 | 98.55413293 |
| RWR | 2 1B | | 0.000506312 | 3.35909e-05 | 0.004 | 87.34188474 | 99.160228484 |
| RWR | 4 1B | | 0.000484066 | 3.78031e-05 | 0.004 | 87.89834315 | 99.05374716 |

- Benchmark works best 1MB Block Size and 2 cores. For block Size 10MB It's closer but not optimal.

- 4

- Generally a CPU fetches a batch of RAM positions at once. It then keeps those in a cache while it's working on this. So if the data in RAM is sequential, the CPU has to wait on such fetching a lot less than if the data it needs to work on is scattered throughout RAM.
- It's not a Quad-core machine so using 4 threads actually hurts performance than 2 threads as a lot of time is spent in context switching.
- Interestingly sequential access for 10MB data 4 threads takes more time than Random Access 4 thread 10 MB . As these are first 2 operations done on the machine looks like cache is not warmed up , so this perticular test case results in anomaly.
- pmbw also gives optimal output for block size of 1MB , but instead of 2 threads , result of 4 thread is better . Which is a deviation from my code result.

4 DISK Benchmarking (Run on Hyperion)

We were asked to Write and Read 10GB data. With various access patterns (RWS, RWR) and various block sizes (1MB, 10MB, 100MB) using 1 thread, 2 threads and 4 threads.

IOZone benchmark

```

root@hyperion:~# ./iozone -s 1024 -t 4 -i 1 -l 2 -r 1024 -w 1024
iozone: Performance Test of File I/O
Version: 3.50a
Compiled for 64 bit arch.
Build Time:

Contributors: William Norcott, Gus Dapp, Ivan Crawford, Kirby Collins,
Al Gorton, Scott Shyne, Mike Warner, Ken Oye,
Steve Laskary, David Smith, Mike Kelly, Ar. Alain Chb,
Markus Luecke, Rene Henning, Ken Hillory, Geoff Haines,
Gordon Sayers, Jeff Kaufman, Wendy Hickey, David Re,
Eric Nishigaki, Mike Carpenter, Michael King, Graham Kells,
Krzysztof Ruchalski, Zhenghui Ren, Qin Li, Gordon Sayers,
Ken Englund.

Run begins: Sun May 20 24:30:50 2018

Auto Mode
File size set to 102400 KB
Record Size: 1024 KB
Record Size: 1024 KB
Command Line Used: ./iozone -s 1024 -t 4 -i 1 -l 2 -r 1024 -w 1024
Output is in /tmp/iozone.
Test Resolution is 0.00000 seconds.
Processor cache size set to 32KB bytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

      random random      read      read      write      write
iozone -s 1024 1024KB 1024KB 30000 32000 25000 25000 25000 25000
iozone -s 1024 1024KB 1024KB 30000 32000 25000 25000 25000 25000

```

corresponding results re kept in /cs553-pa1/disk/iozone.sh

Running Disk Program

```

root@hyperion:~# ./cs553-pa1/disk/iozone.sh
iozone: Performance Test of File I/O
Version: 3.50a
Compiled for 64 bit arch.
Build Time:

Contributors: William Norcott, Gus Dapp, Ivan Crawford, Kirby Collins,
Al Gorton, Scott Shyne, Mike Warner, Ken Oye,
Steve Laskary, David Smith, Mike Kelly, Ar. Alain Chb,
Markus Luecke, Rene Henning, Ken Hillory, Geoff Haines,
Gordon Sayers, Jeff Kaufman, Wendy Hickey, David Re,
Eric Nishigaki, Mike Carpenter, Michael King, Graham Kells,
Krzysztof Ruchalski, Zhenghui Ren, Qin Li, Gordon Sayers,
Ken Englund.

Run begins: Sun May 20 24:30:50 2018

Auto Mode
File size set to 102400 KB
Record Size: 1024 KB
Record Size: 1024 KB
Command Line Used: ./iozone -s 1024 -t 4 -i 1 -l 2 -r 1024 -w 1024
Output is in /tmp/iozone.
Test Resolution is 0.00000 seconds.
Processor cache size set to 32KB bytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

      random random      read      read      write      write
iozone -s 1024 1024KB 1024KB 30000 32000 25000 25000 25000 25000
iozone -s 1024 1024KB 1024KB 30000 32000 25000 25000 25000 25000

```

Disk benchmark Table

| Work- load | Concurrency | Block Size | MyDiskBench Measured Throughput (MB/sec) | IOZone Measured Throughput (MB/sec) | Theoretical Throughput (MB/sec) | MyDiskBench Efficiency (%) | IOZone Efficiency (%) |
|------------|-------------|------------|---|--|------------------------------------|-------------------------------|--------------------------|
| RS | 1 | 1MB | 24.240029 | 35.033 | 1200 | 2.020002417 | 29.84108333 |
| | 1 | 10MB | 22.722776 | 303.734 | 1200 | 1.893564667 | 25.31116667 |
| | 1 | 100MB | 17.949336 | 341.998 | 1200 | 1.495778 | 28.49983333 |
| | 2 | 1MB | 26.432186 | 358.093 | 1200 | 2.202682167 | 29.84108333 |
| | 2 | 10MB | 30.5996 | 303.734 | 1200 | 2.549966667 | 25.31116667 |
| | 2 | 100MB | 31.167425 | 341.998 | 1200 | 2.597285417 | 28.49983333 |
| | 4 | 1MB | 43.434794 | 358.093 | 1200 | 3.619566167 | 29.84108333 |
| | 4 | 10MB | 50.746256 | 303.734 | 1200 | 4.228854667 | 25.31116667 |
| | 4 | 100MB | 38.324509 | 341.998 | 1200 | 3.193709083 | 28.49983333 |
| | 1 | 1MB | 43.196197 | 836.891 | 1200 | 3.599683083 | 69.74091667 |
| | 1 | 10MB | 36.822835 | 688.494 | 1200 | 3.068569583 | 57.3745 |
| | 1 | 100MB | 38.265487 | 546.319 | 1200 | 3.188790583 | 45.52658333 |
| WS | 2 | 1MB | 43.337197 | 836.891 | 1200 | 3.611433083 | 69.74091667 |
| | 2 | 10MB | 45.127668 | 688.494 | 1200 | 3.760639 | 57.3745 |
| | 2 | 100MB | 45.08773 | 546.319 | 1200 | 3.757310833 | 45.52658333 |
| | 4 | 1MB | 48.600181 | 836.891 | 1200 | 4.050015083 | 69.74091667 |
| | 4 | 10MB | 38.434767 | 688.494 | 1200 | 3.20289725 | 57.3745 |
| | 4 | 100MB | 38.099094 | 546.319 | 1200 | 3.1749245 | 45.52658333 |
| | 1 | 1MB | 11230.61736 | 257.997 | 1200 | 935.8847798 | 21.49975 |
| | 1 | 10MB | 3317.251299 | 312.902 | 1200 | 276.4376083 | 26.07516667 |
| | 1 | 100MB | 3015.231274 | 337.239 | 1200 | 251.2692728 | 28.10325 |
| | 2 | 1MB | 1019.43294 | 257.997 | 1200 | 84.952745 | 21.49975 |
| | 2 | 10MB | 382.129398 | 312.902 | 1200 | 31.8441165 | 26.07516667 |
| | 2 | 100MB | 106.100706 | 337.239 | 1200 | 8.8417255 | 28.10325 |
| RR | 4 | 1MB | 7148.759976 | 257.997 | 1200 | 595.729998 | 21.49975 |
| | 4 | 10MB | 5366.675414 | 312.902 | 1200 | 447.2229512 | 26.07516667 |
| | 4 | 100MB | 527.133318 | 337.239 | 1200 | 43.9277765 | 28.10325 |
| | 1 | 1MB | 232.433107 | 908.776 | 1200 | 19.36942558 | 75.73133333 |
| | 1 | 10MB | 850.706823 | 968.512 | 1200 | 70.89223525 | 80.70933333 |
| | 1 | 100MB | 841.7931621 | 851.317 | 1200 | 70.14943018 | 70.94308333 |
| | 2 | 1MB | 270.136348 | 908.776 | 1200 | 22.51136233 | 75.73133333 |
| | 2 | 10MB | 169.15893 | 968.512 | 1200 | 14.0965775 | 80.70933333 |
| | 2 | 100MB | 193.636534 | 851.317 | 1200 | 16.13637783 | 70.94308333 |
| | 4 | 1MB | 244.584262 | 908.776 | 1200 | 20.38202183 | 75.73133333 |
| | 4 | 10MB | 172.471967 | 968.512 | 1200 | 14.37266392 | 80.70933333 |
| | 4 | 100MB | 235.074646 | 851.317 | 1200 | 19.58955383 | 70.94308333 |

Synopsis for Disk Benchmark

theoretical throughput = one lane rate * 2 lane per port * 1 port per stack
= 600 MB/sec * 2 * 1 = 1200 MB/sec.

- Divided all 36 different disk throughput operations in 9 groups and parallelly run them on 9 nodes.
- Disk gives highest throughput when in Sequential Read , using 1 thread and read in 100 MB chunk.
- Creating Disk Write first followed by Read so no need to copy temp files .
-

5 Network Benchmarking

We were asked to implement Server-Client software using both TCP and UDP protocols with Workload: 1GB data. We operated over it 100X times with various block sizes (1KB, 32KB) with Concurrency: 1 thread, 2 threads, 4 threads and 8 threads.

Theoretical network bandwidth calculation

Theoretical Throughput = $\frac{RWIN}{RTT}$

Where RWIN is the TCP Receive Window and RTT is the round-trip time for the path. The Max TCP Window size in the absence of TCP window scale option is 65,535 bytes.

Iperf results UDP

```
schatterjee@compute-4: ~ (ssh)
schatterjee@compute-4: ~$ iperf -c 172.16.1.4 --port 4030 --format m --len 32K --udp -i 5 -b 1703M
Client connecting to 172.16.1.4, UDP port 4030
Sending 32768 byte datagrams
UDP buffer size: 0.20 MByte (default)

[ 3] local 172.16.1.4 port 4744t connected with 172.16.1.4 port 4030
ID Interval Transfer Bandwidth
[ 3] 0.0- 5.0 sec 1015 Mbytes 1702 Mbits/sec
[ 3] 5.0-10.0 sec 1015 Mbytes 1702 Mbits/sec
[ 3] 0.0-10.0 sec 2032 Mbytes 1702 Mbits/sec
[ 3] Sent 64006 datagrams
[ 3] Server Report:
[ 3] 0.0-10.0 sec 2018 Mbytes 1693 Mbits/sec 0.040 ms 303/64005 (0.50%)
[ 3] 0.0-10.0 sec 1 datagrams received out-of-order
```

```
schatterjee@prometheus: ~ (ssh)
schatterjee@prometheus: ~$ iperf -s --port 4030 --format m --len 32K --udp
Server listening on UDP port 4030
Receiving 32768 byte datagrams
UDP buffer size: 0.20 MByte (default)

[ 3] local 172.16.1.4 port 4030 connected with 172.16.1.4 port 4744t
ID Interval Transfer Bandwidth Jitter Lost/Total Datagrams
[ 3] 0.0-10.0 sec 2018 Mbytes 1693 Mbits/sec 0.041 ms 303/64005 (0.50%)
[ 3] 0.0-10.0 sec 1 datagrams received out-of-order
```

Iperf results TCP

```
schatterjee@compute-9: ~$ iperf -c 172.16.1.4 --port 4030 --format m --len 32K --run 10240M -i 10
Client connecting to 172.16.1.4, TCP port 4030
TCP window size: 0.04 MByte (default)

[ 3] local 172.16.1.4 port 33930 connected with 172.16.1.4 port 4030
ID Interval Transfer Bandwidth
[ 3] 0.0-10.0 sec 4710 Mbytes 3551 Mbits/sec
[ 3] 10.0-20.0 sec 4742 Mbytes 3978 Mbits/sec
[ 3] 0.0-21.6 sec 10240 Mbytes 3972 Mbits/sec
schatterjee@compute-9: ~$
```

```
4] local 172.16.1.4 port 4030 connected with 172.16.1.4 port 33930
ID Interval Transfer Bandwidth
4] 0.0-10.0 sec 1307 Mbytes 1004 Mbits/sec
5] local 172.16.1.4 port 4030 connected with 172.16.1.4 port 33910
5] 0.0-0.0 sec 9853 Mbytes 1215 Mbits/sec
4] local 172.16.1.4 port 4030 connected with 172.16.1.4 port 33912
4] 0.0-72.2 sec 10240 Mbytes 4190 Mbits/sec
5] local 172.16.1.4 port 4030 connected with 172.16.1.4 port 33920
5] 0.0-21.6 sec 10240 Mbytes 3970 Mbits/sec
```

Rest of the Iperf commands are detailed in iperf.sh

Running TCP BenchMark Codes on Hyperion

```
./scripts/home/schatterjee/iperf3-pat/network/MACBench-TCP.app: In function 'int main(int, char**):
./scripts/home/schatterjee/iperf3-pat/network/MACBench-TCP.app:105:60: warning: ISO C++ forbids converting a string constant to 'char*' [-Werror=stringop-overflow]
    if (storeInfo(customFile.c_str(), hostName, "w")) out << "0
./scripts/home/schatterjee/iperf3-pat/network/MACBench-TCP.app:205:63: warning: integer overflow in expression [-Werror=overflow]
    * throughput in Mps: " + toString(100 * oneGBdiv
./scripts/home/schatterjee/iperf3-pat/network/MACBench-TCP.app:210:43: warning: ISO C++ forbids converting a string constant to 'char*' [-Werror=stringop-overflow]
    storeInfo(reportPath, report.c_str());
./scripts/home/schatterjee/iperf3-pat/network/MACBench-TCP.app: In function 'void readAndWriteFromState(char*)':
./scripts/home/schatterjee/iperf3-pat/network/MACBench-TCP.app:52:20: warning: ignoring return value of 'int ReadFile(LPCSTR, LPVOID, DWORD, FILE*, void*)' declared with attribute 'warn_unused_result' (unused-result)
    ReadFile(file, "a", buff);
Folder /allum created.
Submitted batch job 23454
Submitted batch job 23455
Submitted batch job 23456
Submitted batch job 23457
Submitted batch job 23458
Submitted batch job 23459
Submitted batch job 23460
Submitted batch job 23461
Submitted batch job 23462
schatterjee@hyperion: ~$
```

```
===== Time: Tue Mar 27 12:18:13 CDT 2018. New Execution =====
Threads 2, BLOCKSIZE -3000, TotalTime = 1.06229, Node client, Total Data Received 1257664000, Throughput in Mps: 455.8439, Server Node redcompute-3
Threads 2, BLOCKSIZE -3000, TotalTime = 1.81492, Node server, Total Data Received 1257664000, Throughput in Mps: 7.94678, Server Node redcompute-5
Threads 4, BLOCKSIZE -3000, TotalTime = 1.04939, Node client, Total Data Received 1257664000, Throughput in Mps: 5813.62487, Server Node redcompute-3
Threads 4, BLOCKSIZE -3000, TotalTime = 157.77000, Node server, Total Data Received 1257664000, Throughput in Mps: 61.60329, Server Node redcompute-3
Threads 8, BLOCKSIZE -3000, TotalTime = 1.03091, Node client, Total Data Received 1257664000, Throughput in Mps: 5848.40554, Server Node redcompute-2
Threads 8, BLOCKSIZE -3000, TotalTime = 459.02432, Node server, Total Data Received 1257664000, Throughput in Mps: 60.770514, Server Node redcompute-2
Threads 1, BLOCKSIZE -3000, TotalTime = 6.616527, Node client, Total Data Received 1257664000, Throughput in Mps: 4470.443010, Server Node redcompute-6
Threads 1, BLOCKSIZE -3000, TotalTime = 422.340048, Node server, Total Data Received 1257664000, Throughput in Mps: 10.207056, Server Node redcompute-6
schatterjee@hyperion: ~$
```

```
1 redcompute-6 2944 compute netClique schatter R 6.47
1 redcompute-14 2942 compute netClique schatter R 6.47
1 redcompute-2 2942 compute netClique schatter R 6.47
1 redcompute-3 2945 compute netClique schatter R 6.47
2 redcompute-2 2947 compute UP-4-10 sajnera4 R 1.04
2 redcompute-1 2945 interacti bash dsalman R 0.14
1 bilocompute-2 2945 interacti bash dsalman R 0.14
schatterjee@hyperion: ~$ square
8 NODE1351(64520) 2233 compute MyGiside cyrusale CD 15.26
1 redcompute-1 23415 interacti bash dsalman R 11.59
1 bilocompute-1 23420 interacti bash dsalman R 20.59
1 bilocompute-3 2947 compute UP-4-10 sajnera4 R 6.24
2 redcompute-1 2945 interacti bash dsalman R 5.31
1 bilocompute-2 2945 compute bash dsalman R 3.27
4 redcompute-4 2959 compute bash dsalman R 3.23
4 redcompute-4 2959 compute UP-4-10 sajnera4 R 0.39
2 redcompute-1 2963 compute UP-8-10 sajnera4 R 0.36
schatterjee@hyperion: ~$ square
```

Throughput calculation Table

| Protocol | Concurrency | Block Size | MyNETBench Measured Throughput (Mb/sec) | iperf Measured Throughput (Mb/sec) | Theoretical Throughput (Mb/sec) | MyNETBench Efficiency (%) | iperf Efficiency (%) |
|----------|-------------|------------|---|------------------------------------|---------------------------------|---------------------------|----------------------|
| TCP | 1 | 1KB | 1.2389 | 564 | 1323.939394 | 0.09357679 | 42.6001373 |
| TCP | 1 | 32KB | 1.6784 | 3626 | 1137.266811 | 0.1475819 | 318.834592 |
| TCP | 2 | 1KB | 1.2389 | 566 | 1323.939394 | 0.09357679 | 42.7512016 |
| TCP | 2 | 32KB | 1.6784 | 3666 | 1137.266811 | 0.1475819 | 322.351797 |
| TCP | 4 | 1KB | 1.2389 | 562 | 1323.939394 | 0.09357679 | 42.449073 |
| TCP | 4 | 32KB | 1.6784 | 3,524.00 | 1137.266811 | 0.1475819 | 309.865721 |
| TCP | 8 | 1KB | 1.2389 | 563 | 1323.939394 | 0.09357679 | 42.5246052 |
| TCP | 8 | 32KB | 1.6784 | 3585 | 1137.266811 | 0.1475819 | 315.229458 |
| UDP | 1 | 1KB | 0.9765 | 559 | 1323.939394 | 0.07375715 | 42.2224765 |
| UDP | 1 | 32KB | 0.9218 | 1693 | 1137.266811 | 0.08105398 | 148.865682 |
| UDP | 2 | 1KB | 0.9765 | 892 | 1323.939394 | 0.07375715 | 67.3746853 |
| UDP | 2 | 32KB | 0.9218 | 1761 | 1137.266811 | 0.08105398 | 154.84493 |
| UDP | 4 | 1KB | 0.9765 | 987 | 1323.939394 | 0.07375715 | 74.5502403 |
| UDP | 4 | 32KB | 0.9218 | 1198 | 1137.266811 | 0.08105398 | 105.340276 |
| UDP | 8 | 1KB | 0.9765 | 959 | 1323.939394 | 0.07375715 | 72.4353399 |
| UDP | 8 | 32KB | 0.9218 | 1399 | 1137.266811 | 0.08105398 | 123.014229 |

Synopsis for bandwidth calculation

- Increasing No of thread does have significant +ve effect on Measured Throughput . Looks like 1 thread cant max out TCP buffer , If we use 8 threads then TCP throughput is maximal . For UDP it's 2 threads
- for TCP best result is obtained when block size is 32kb and thread count is 8.
- TCP is giving effectively higher Throughput than UDP . Each frame goes through several buffers we send it: The application buffer, The Protocol Buffer, The Software interface buffer and the Hardware interface buffer. As we start stressing the stack by sending high speed data you will fill up these buffers and TCP is performing better as TCP is optimized for high speed bulk transfers.
- Throughput is dependent on Block size as higher block size yeilds in greater utilization of the available bandwidth and less no of iteration.

Latency calculation using Pingpong.cpp and 'PING' utility tool

Implemented similar Server-Client software to check latency, where server and client need to acknowledge the receipt of 1B data. And iterated it 1 million times.

Theoretical Latency: Ultimately response time over a network is limited by the speed of light. In a vacuum, light travels with a speed of 299 792 458 m / s. So, if we assume that a 'ping' travels with the speed of light, which are the best possible response times we can get ? Assuming machines are 10m apart . Latency should be $\frac{10*2*1000}{299792458} \text{ MilliSecond} = 0.000066ms$

Executing PingPong.cpp

```
schatterjee@hyperionides:~$ ./cs553-pat/pingpong_engine.sh 1
/exports/home/schatterjee/cs553-pat/pingpong/Pingpong.cpp:36:72: warning: ISO C++ forbids converting a string constant to 'char*' [-Wwrite-strings]
int storeInFile(const char *filename, const char *data, char *mode="a+") {
/exports/home/schatterjee/cs553-pat/pingpong/Pingpong.cpp: In function 'int main(int, char**)':
/exports/home/schatterjee/cs553-pat/pingpong/Pingpong.cpp:235:68: warning: ISO C++ forbids converting a string constant to 'char*' [-Wwrite-strings]
if (!storeInFile(customfile.c_str(), host_name, "w")) out << "Couldn't write IP to
/exports/home/schatterjee/cs553-pat/pingpong/Pingpong.cpp:285:43: warning: ISO C++ forbids converting a string constant to 'char*' [-Wwrite-strings]
storeInFile(reportPath, report.c_str());
/exports/home/schatterjee/cs553-pat/pingpong/Pingpong.cpp: In function 'void readFromFile(const char*, char*)':
/exports/home/schatterjee/cs553-pat/pingpong/Pingpong.cpp:59:29: warning: ignoring return value of 'int fscanf(FILE*, const char*, ...)', declared with attribute warn_unused_result [-Wunused-result]
fscanf(file, "%s", buff);
Folder /slurm created..
Submitted batch job 27345
Submitted batch job 27346
Submitted batch job 27347
Submitted batch job 27348
Submitted batch job 27349
Submitted batch job 27351
Submitted batch job 27352
Submitted batch job 27353
schatterjee@hyperionides:~$ {}
```

```
=> ./slurms/pingpongserver.out <=
Threads 4, TotalTime = 235.005263, Mode server, Total Data Received 1000002, Latency in ms: 0.233005, Server Node redcompu
te-2, Transmission Mode: TCP..

schatterjee@hyperionides:~$ ssh

===== Time: Wed Mar 28 21:08:56 CDT 2018. New Execution =====
Threads 2, TotalTime = 200.66376, Mode server, Total Data Received 1000001, Latency in ms: 0.200664, Server Node redcompu
te-3, Transmission Mode: TCP..
Threads 8, TotalTime = 38.276290, Mode client, Total Data Received 1000005, Latency in ms: 0.039276, Server Node redcompu
te-4, Transmission Mode: TCP..
Threads 8, TotalTime = 224.411917, Mode server, Total Data Received 1000005, Latency in ms: 0.224412, Server Node redcompu
te-4, Transmission Mode: TCP..
Threads 1, TotalTime = 162.873105, Mode client, Total Data Received 1000000, Latency in ms: 0.162873, Server Node redcompu
te-5, Transmission Mode: TCP..
Threads 1, TotalTime = 222.120745, Mode server, Total Data Received 1000000, Latency in ms: 0.222121, Server Node redcompu
te-5, Transmission Mode: TCP..
Threads 4, TotalTime = 235.005263, Mode server, Total Data Received 1000002, Latency in ms: 0.233005, Server Node redcompu
te-2, Transmission Mode: TCP..
```

Latency calculation Table

| Protocol | Concurrency | Message Size | MyNETBench Measured Latency (ms) | ping Measured Latency (ms) | Theoretical Latency (ms) | MyNETBench Efficiency (%) | Ping Efficiency |
|----------|-------------|--------------|----------------------------------|----------------------------|--------------------------|---------------------------|-----------------|
| TCP | 1 | 1B | 0.186475 | 0.00396 | 0.000066 | 0.035393484 | 1.6666667 |
| TCP | 2 | 1B | 0.100287 | 0.00396 | 0.000066 | 0.065811122 | 1.6666667 |
| TCP | 4 | 1B | 0.168867 | 0.00396 | 0.000066 | 0.039084013 | 1.6666667 |
| TCP | 8 | 1B | 0.224412 | 0.00396 | 0.000066 | 0.029410192 | 1.6666667 |
| UDP | 1 | 1B | 0.051648 | 0.00396 | 0.000066 | 0.127788104 | 1.6666667 |
| UDP | 2 | 1B | 0.163044 | 0.00396 | 0.000066 | 0.04047987 | 1.6666667 |
| UDP | 4 | 1B | 0.107404 | 0.00396 | 0.000066 | 0.061450225 | 1.6666667 |
| UDP | 8 | 1B | 0.057742 | 0.00396 | 0.000066 | 0.114301548 | 1.6666667 |

Synopsis for Latency calculation

- UDP is faster than TCP, and the simple reason is because its nonexistent acknowledge packet (ACK) that permits a continuous packet stream, instead of TCP that acknowledges a set of packets, calculated by using the TCP window size and round-trip time (RTT).
- UDP latencies are less compared to TCP
- For TCP we get best latency with 2 threads. And for UDP it's 1 thread.

References

- [1] Measuring network throughput
[https : //en.wikipedia.org/wiki/Measuring_network_throughput](https://en.wikipedia.org/wiki/Measuring_network_throughput)
- [2] UDP Latency
[https : //stackoverflow.com/questions/47903/udp-vs-tcp-how-much-faster-is-it](https://stackoverflow.com/questions/47903/udp-vs-tcp-how-much-faster-is-it)
- [3] TCP Bandwidth
[https : //serverfault.com/questions/432101/why-is-udp-slower-than-tcp-on-ubuntu-server](https://serverfault.com/questions/432101/why-is-udp-slower-than-tcp-on-ubuntu-server)